

# dog\_app

September 10, 2019

## 1 Convolutional Neural Networks

### 1.1 Project: Write an Algorithm for a Dog Identification App

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

## Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

**Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.**

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog\_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

*Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.*

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human\_files and dog\_files.

```
In [40]: import numpy as np
         from glob import glob

         # load filenames for human and dog images
         human_files = np.array(glob("/data/lfw/*/"))
         dog_files = np.array(glob("/data/dog_images/*/"))

         # print number of images in each dataset
         print('There are %d total human images.' % len(human_files))
         print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

### ## Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [41]: import cv2
         import matplotlib.pyplot as plt
         %matplotlib inline

         # extract pre-trained face detector
         face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

         # load color (BGR) image
         img = cv2.imread(human_files[0])
         # convert BGR image to grayscale
         gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

         # find faces in image
         faces = face_cascade.detectMultiScale(gray)

         # print number of faces detected in the image
         print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

### 1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [42]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

### 1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

**Question 1:** Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

**Answer:**

- Percentage of first 100 images in `human_files` which detected a human face: 98%
- Percentage of first 100 images in `dog_files` which detected a human face: 17%

```
In [43]: from tqdm import tqdm
```

```
human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.

human_true = sum(1 for file in human_files_short if face_detector(file))
dog_false = sum(1 for file in dog_files_short if face_detector(file))

print("Percentage of human faces correctly detected as human face: " + str(human_true/100))
print("Percentage of dog faces incorrectly detected as human face: " + str(dog_false/100))
```

```
Percentage of human faces correctly detected as human face: 0.98
```

```
Percentage of dog faces incorrectly detected as human face: 0.17
```

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make

use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)  
       ### TODO: Test performance of another face detection algorithm.  
       ### Feel free to use as many code cells as needed.
```

---

## ## Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

### 1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [44]: import torch  
        import torchvision.models as models  
  
        # define VGG16 model  
        VGG16 = models.vgg16(pretrained=True)  
  
        # check if CUDA is available  
        use_cuda = torch.cuda.is_available()  
  
        # move model to GPU if CUDA is available  
        if use_cuda:  
            VGG16 = VGG16.cuda()
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

### 1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as `'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg'`) as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [45]: from PIL import Image  
        import torchvision.transforms as transforms  
  
        def VGG16_predict(img_path):
```

```

'''
Use pre-trained VGG-16 model to obtain index corresponding to
predicted ImageNet class for image at specified path

Args:
    img_path: path to an image

Returns:
    Index corresponding to VGG-16 model's prediction
'''

## TODO: Complete the function.
## Load and pre-process an image from the given img_path
## Return the *index* of the predicted class for that image
img = Image.open(img_path)
transform_pipeline = transforms.Compose([transforms.Resize((224,224)),
                                         transforms.ToTensor(),
                                         transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                                std=[0.229, 0.224, 0.225])])

img = transform_pipeline(img)
img = img.unsqueeze(0)
img = img.to('cuda')
VGG16.eval()
prediction = VGG16(img)
prediction = prediction.to('cpu')
prediction = prediction.data.numpy().argmax()
return prediction # predicted class index

```

### 1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns True if a dog is detected in an image (and False if not).

```

In [46]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    if VGG16_predict(img_path) >= 151 and VGG16_predict(img_path) <= 268:
        return True
    else:
        return False

```

### 1.1.6 (IMPLEMENTATION) Assess the Dog Detector

**Question 2:** Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

**Answer:** - Percentage of images in `human_files_short` which detected a dog: 0% - Percentage of images in `dog_files_short` which detected a dog: 100%

```
In [47]: ### TODO: Test the performance of the dog_detector function
        ### on the images in human_files_short and dog_files_short
        human_files_detected_dog = sum(1 for file in human_files_short if dog_detector(file))
        dog_files_detected_dog = sum(1 for file in dog_files_short if dog_detector(file))
        print("Percentage of images in human_files_short which detected dog: " +str(human_files_detected_dog/len(human_files_short)))
        print("Percentage of images in dog_files_short which detected dog: " +str(dog_files_detected_dog/len(dog_files_short)))
```

Percentage of images in `human_files_short` which detected dog: 0.0

Percentage of images in `dog_files_short` which detected dog: 1.0

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

---

#### ## Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

---

Brittany	Welsh Springer Spaniel
----------	------------------------

---

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

---

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

---

---

---

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

---

---

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

---

---

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

---

---

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

### 1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [9]: data_dir = '/data/dog_images'
        train_dir = data_dir + '/train'
        valid_dir = data_dir + '/valid'
        test_dir = data_dir + '/test'

In [10]: import os
         from torchvision import datasets, transforms

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes
         train_transforms = transforms.Compose([transforms.Resize(256),
                                                transforms.CenterCrop(224),
                                                transforms.RandomRotation(30),
                                                transforms.RandomHorizontalFlip(),
                                                transforms.ToTensor(),
                                                transforms.Normalize([0.485, 0.456, 0.406],
                                                                    [0.229, 0.224, 0.225])])

         valid_transforms = transforms.Compose([transforms.Resize(256),
                                                transforms.CenterCrop(224),
                                                transforms.ToTensor(),
                                                transforms.Normalize([0.485, 0.456, 0.406],
                                                                    [0.229, 0.224, 0.225])])
```



```

test_transforms = transforms.Compose([transforms.Resize(256),
                                     transforms.CenterCrop(224),
                                     transforms.ToTensor(),
                                     transforms.Normalize([0.485, 0.456, 0.406],
                                                         [0.229, 0.224, 0.225])])

# TODO: Load the datasets with ImageFolder
train_data = datasets.ImageFolder(train_dir, transform = train_transforms)
valid_data = datasets.ImageFolder(valid_dir, transform = valid_transforms)
test_data = datasets.ImageFolder(test_dir, transform = test_transforms)

# TODO: Using the image datasets and the trainforms, define the dataloaders
trainloader = torch.utils.data.DataLoader(train_data, batch_size=32, shuffle=True)
validloader = torch.utils.data.DataLoader(valid_data, batch_size=32)
testloader = torch.utils.data.DataLoader(test_data, batch_size=32)

loaders_scratch = {'train': trainloader, 'valid': validloader, 'test': testloader}

```

**Question 3:** Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

**Answer:** - My code resizes all images to 256x256 pixels and then center crops them to reduce them to 224x224 pixels, which is the size of the input tensor. This transformation is applied to all train, validation and test datasets - For training dataset in addition to resizing and cropping, I also used additional transformations like random horizontal flips and random rotation by 30 degrees. The aim was to augment the training dataset, so This hat it generalizes better and gives better accuracy on the testing dataset.

### 1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [11]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        self.conv1 = nn.Conv2d(3, 16, 3, padding=1)
        # convolutional layer
        self.conv2 = nn.Conv2d(16, 32, 3, padding=1)
        # convolutional layer
        self.conv3 = nn.Conv2d(32, 64, 3, padding=1)
        # convolutional layer

```

```

self.conv4 = nn.Conv2d(64, 128, 3, padding=1)
# max pooling layer
self.pool = nn.MaxPool2d(2, 2)
# linear layer
self.fc1 = nn.Linear(14 * 14 * 128, 1000)
# linear layer
self.fc2 = nn.Linear(1000, 500)
# linear layer
self.fc3 = nn.Linear(500, 133)
# dropout layer (p=0.20)
self.dropout = nn.Dropout(0.20)

def forward(self, x):
    ## Define forward behavior
    x = self.pool(F.relu(self.conv1(x)))
    x = self.pool(F.relu(self.conv2(x)))
    x = self.pool(F.relu(self.conv3(x)))
    x = self.pool(F.relu(self.conv4(x)))
    # flatten image input
    x = x.view(-1, 128 * 14 * 14)
    # add dropout layer
    x = self.dropout(x)
    # add 1st hidden layer, with relu activation function
    x = F.relu(self.fc1(x))
    # add dropout layer
    x = self.dropout(x)
    # add 2nd hidden layer, with relu activation function
    x = F.relu(self.fc2(x))
    # add dropout layer
    x = self.dropout(x)
    x = self.fc3(x)
    return x

### You so NOT have to modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()
print(model_scratch)

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

```

Net(
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

(pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(fc1): Linear(in_features=25088, out_features=1000, bias=True)
(fc2): Linear(in_features=1000, out_features=500, bias=True)
(fc3): Linear(in_features=500, out_features=133, bias=True)
(dropout): Dropout(p=0.2)
)

```

**Question 4:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

**Answer:** - My final CNN architecture consists of 4 Convolutional layers and MaxPooling layers followed by 3 Fully connected layers. I have used 16 3x3 filters with padding of 1 for the first convolutional layer and gradually increasing the number of filters by a factor of 2. For Max Pooling layers I reduce the length and breadth by a factor of 2 where as increasing the depth. Finally we have 3 fully connected layers with 1000, 500 and 133 (number of classes) units. I have used RELU activation function for the layers and have used Dropout regularization with  $p=0.2$  between the FC layers. - I initially used 3 Convolutional layers and 2 fully connected layers but I faced issues to meeting the 10% test accuracy rubric as specified by the project. So I decided to add one more convolutional layer and 1 more fully connected layer.

### 1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [12]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.Adam(model_scratch.parameters(), lr=0.001)

In [13]: from PIL import ImageFile
        ImageFile.LOAD_TRUNCATED_IMAGES = True

```

### 1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```

In [14]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""
        # initialize tracker for minimum validation loss
        valid_loss_min = np.Inf

        for epoch in range(1, n_epochs+1):
            # initialize variables to monitor training and validation loss
            train_loss = 0.0

```

```

valid_loss = 0.0

#####
# train the model #
#####
model.train()
for batch_idx, (data, target) in enumerate(loaders['train']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## find the loss and update the model parameters accordingly
    ## record the average training loss, using something like
    ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
    optimizer.zero_grad()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the batch loss
    loss = criterion(output, target)
    # backward pass: compute gradient of the loss with respect to model parameters
    loss.backward()
    # perform a single optimization step (parameter update)
    optimizer.step()
    # update training loss
    train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))

#####
# validate the model #
#####
model.eval()
with torch.no_grad():
    for batch_idx, (data, target) in enumerate(loaders['valid']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## update the average validation loss
        output = model(data)
        # calculate the batch loss
        loss = criterion(output, target)
        # update average validation loss
        valid_loss += ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(

## TODO: save the model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...

```

```

        torch.save(model.state_dict(), save_path)
        valid_loss_min = valid_loss
    # return trained model
    return model

# train the model
model_scratch = train(30, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

```

Epoch: 1      Training Loss: 4.850713      Validation Loss: 4.725215
Validation loss decreased (inf --> 4.725215). Saving model ...
Epoch: 2      Training Loss: 4.632405      Validation Loss: 4.556520
Validation loss decreased (4.725215 --> 4.556520). Saving model ...
Epoch: 3      Training Loss: 4.527247      Validation Loss: 4.472876
Validation loss decreased (4.556520 --> 4.472876). Saving model ...
Epoch: 4      Training Loss: 4.370166      Validation Loss: 4.247534
Validation loss decreased (4.472876 --> 4.247534). Saving model ...
Epoch: 5      Training Loss: 4.226288      Validation Loss: 4.160502
Validation loss decreased (4.247534 --> 4.160502). Saving model ...
Epoch: 6      Training Loss: 4.122058      Validation Loss: 4.118450
Validation loss decreased (4.160502 --> 4.118450). Saving model ...
Epoch: 7      Training Loss: 4.028961      Validation Loss: 4.059320
Validation loss decreased (4.118450 --> 4.059320). Saving model ...
Epoch: 8      Training Loss: 3.911133      Validation Loss: 3.920666
Validation loss decreased (4.059320 --> 3.920666). Saving model ...
Epoch: 9      Training Loss: 3.802415      Validation Loss: 3.943581
Epoch: 10     Training Loss: 3.706408      Validation Loss: 3.876677
Validation loss decreased (3.920666 --> 3.876677). Saving model ...
Epoch: 11     Training Loss: 3.587921      Validation Loss: 3.834834
Validation loss decreased (3.876677 --> 3.834834). Saving model ...
Epoch: 12     Training Loss: 3.484577      Validation Loss: 3.705796
Validation loss decreased (3.834834 --> 3.705796). Saving model ...
Epoch: 13     Training Loss: 3.393525      Validation Loss: 3.729004
Epoch: 14     Training Loss: 3.292607      Validation Loss: 3.661549
Validation loss decreased (3.705796 --> 3.661549). Saving model ...
Epoch: 15     Training Loss: 3.217931      Validation Loss: 3.598773
Validation loss decreased (3.661549 --> 3.598773). Saving model ...
Epoch: 16     Training Loss: 3.143845      Validation Loss: 3.577302
Validation loss decreased (3.598773 --> 3.577302). Saving model ...
Epoch: 17     Training Loss: 3.025565      Validation Loss: 3.698697
Epoch: 18     Training Loss: 2.963040      Validation Loss: 3.679426
Epoch: 19     Training Loss: 2.869311      Validation Loss: 3.659032
Epoch: 20     Training Loss: 2.798219      Validation Loss: 3.718408
Epoch: 21     Training Loss: 2.714042      Validation Loss: 3.684733

```

Epoch: 22	Training Loss: 2.656363	Validation Loss: 3.716555
Epoch: 23	Training Loss: 2.574896	Validation Loss: 3.536162
Validation loss decreased (3.577302 --> 3.536162). Saving model ...		
Epoch: 24	Training Loss: 2.506106	Validation Loss: 3.652049
Epoch: 25	Training Loss: 2.436637	Validation Loss: 3.546527
Epoch: 26	Training Loss: 2.349528	Validation Loss: 3.752162
Epoch: 27	Training Loss: 2.287557	Validation Loss: 3.650603
Epoch: 28	Training Loss: 2.244469	Validation Loss: 3.769622
Epoch: 29	Training Loss: 2.183292	Validation Loss: 3.777796
Epoch: 30	Training Loss: 2.119282	Validation Loss: 3.851401

### 1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [15]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    with torch.no_grad():
        for batch_idx, (data, target) in enumerate(loaders['test']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            # forward pass: compute predicted outputs by passing inputs to the model
            output = model(data)
            # calculate the loss
            loss = criterion(output, target)
            # update average test loss
            test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
            # convert output probabilities to predicted class
            pred = output.data.max(1, keepdim=True)[1]
            # compare predictions to true label
            correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
            total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (100. * correct / total, correct, total))

    # call test function
    test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

Test Loss: 3.581942

Test Accuracy: 19% (167/836)

## 2 Accuracy of the CNN model built from scratch

- The CNN built from scratch achieves a test accuracy of 19% on the test dataset which exceeds the 10% test accuracy rubric specified for the project

---

## Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

### 2.0.1 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [22]: data_dir = '/data/dog_images'
        train_dir = data_dir + '/train'
        valid_dir = data_dir + '/valid'
        test_dir = data_dir + '/test'

In [23]: import os
        import torch
        from torchvision import datasets, transforms

In [24]: ## TODO: Specify data loaders
        train_transforms = transforms.Compose([transforms.Resize(256),
                                              transforms.CenterCrop(224),
                                              transforms.RandomRotation(30),
                                              transforms.RandomHorizontalFlip(),
                                              transforms.ToTensor(),
                                              transforms.Normalize([0.485, 0.456, 0.406],
                                                                  [0.229, 0.224, 0.225])])

        valid_transforms = transforms.Compose([transforms.Resize(256),
                                              transforms.CenterCrop(224),
                                              transforms.ToTensor(),
                                              transforms.Normalize([0.485, 0.456, 0.406],
                                                                  [0.229, 0.224, 0.225])])
```

```

test_transforms = transforms.Compose([transforms.Resize(256),
                                     transforms.CenterCrop(224),
                                     transforms.ToTensor(),
                                     transforms.Normalize([0.485, 0.456, 0.406],
                                                         [0.229, 0.224, 0.225])])

# TODO: Load the datasets with ImageFolder
train_data = datasets.ImageFolder(train_dir, transform = train_transforms)
valid_data = datasets.ImageFolder(valid_dir, transform = valid_transforms)
test_data = datasets.ImageFolder(test_dir, transform = test_transforms)

# TODO: Using the image datasets and the trainforms, define the dataloaders
trainloader = torch.utils.data.DataLoader(train_data, batch_size=32, shuffle=True)
validloader = torch.utils.data.DataLoader(valid_data, batch_size=32)
testloader = torch.utils.data.DataLoader(test_data, batch_size=32)

loaders_transfer = {'train': trainloader, 'valid': validloader, 'test': testloader}

```

## 2.0.2 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [25]: import torchvision.models as models
import torch.nn as nn

use_cuda = torch.cuda.is_available()

## TODO: Specify model architecture
model_transfer = models.vgg16(pretrained=True)

# Freeze training for all "features" layers
for param in model_transfer.features.parameters():
    param.requires_grad = False

n_inputs = model_transfer.classifier[6].in_features

# add last linear layer (n_inputs -> 5 flower classes)
# new layers automatically have requires_grad = True
last_layer = nn.Linear(n_inputs, 133)

model_transfer.classifier[6] = last_layer
print(model_transfer)

if use_cuda:
    model_transfer = model_transfer.cuda()

```



```

VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace)
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): ReLU(inplace)
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): ReLU(inplace)
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (22): ReLU(inplace)
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (25): ReLU(inplace)
    (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (27): ReLU(inplace)
    (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (29): ReLU(inplace)
    (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Linear(in_features=25088, out_features=4096, bias=True)
    (1): ReLU(inplace)
    (2): Dropout(p=0.5)
    (3): Linear(in_features=4096, out_features=4096, bias=True)
    (4): ReLU(inplace)
    (5): Dropout(p=0.5)
    (6): Linear(in_features=4096, out_features=133, bias=True)
  )
)

```

**Question 5:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

**Answer:** - For transfer learning exercise I started with a VGG16 baseline model. The VGG16 baseline model has 1000 output units where as in our problem we have 133 classes. So I modified the classifier section of the VGG16 to change the number of nodes in the output layer from 1000 to 133. I also froze the parameters for all layers except the classifier part of the network. - Although the network I chose was rather simplistic, I was able to achieve a test classification accuracy of 62% compared to 60% rubric.

### 2.0.3 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [26]: import torch.optim as optim
```

```
In [27]: criterion_transfer = nn.CrossEntropyLoss()
optimizer_transfer = optim.Adam(model_transfer.classifier.parameters(), lr=0.001)
```

### 2.0.4 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [28]: from PIL import ImageFile
```

```
ImageFile.LOAD_TRUNCATED_IMAGES = True
```

```
In [29]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
            optimizer.zero_grad()
            # forward pass: compute predicted outputs by passing inputs to the model
            output = model(data)
```

```

        # calculate the batch loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model parameters
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer.step()
        # update training loss
        train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))

#####
# validate the model #
#####
model.eval()
with torch.no_grad():
    for batch_idx, (data, target) in enumerate(loaders['valid']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## update the average validation loss
        output = model(data)
        # calculate the batch loss
        loss = criterion(output, target)
        # update average validation loss
        valid_loss += ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(

## TODO: save the model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'
    torch.save(model.state_dict(), save_path)
    valid_loss_min = valid_loss

# return trained model
return model

```

In [30]: # train the model

```

model_transfer = train(20, loaders_transfer, model_transfer, optimizer_transfer, criterion_transfer)

# load the model that got the best validation accuracy (uncomment the line below)
model_transfer.load_state_dict(torch.load('model_transfer.pt'))

```

```

Epoch: 1          Training Loss: 4.347772          Validation Loss: 2.299619
Validation loss decreased (inf --> 2.299619). Saving model ...
Epoch: 2          Training Loss: 3.573876          Validation Loss: 1.684224
Validation loss decreased (2.299619 --> 1.684224). Saving model ...
Epoch: 3          Training Loss: 3.205577          Validation Loss: 1.531882

```

```

Validation loss decreased (1.684224 --> 1.531882). Saving model ...
Epoch: 4      Training Loss: 2.985179      Validation Loss: 1.501359
Validation loss decreased (1.531882 --> 1.501359). Saving model ...
Epoch: 5      Training Loss: 2.941744      Validation Loss: 1.416442
Validation loss decreased (1.501359 --> 1.416442). Saving model ...
Epoch: 6      Training Loss: 2.914958      Validation Loss: 1.456971
Epoch: 7      Training Loss: 2.794449      Validation Loss: 1.381727
Validation loss decreased (1.416442 --> 1.381727). Saving model ...
Epoch: 8      Training Loss: 2.731120      Validation Loss: 1.342018
Validation loss decreased (1.381727 --> 1.342018). Saving model ...
Epoch: 9      Training Loss: 2.691458      Validation Loss: 1.339883
Validation loss decreased (1.342018 --> 1.339883). Saving model ...
Epoch: 10     Training Loss: 2.655596      Validation Loss: 1.351768
Epoch: 11     Training Loss: 2.738938      Validation Loss: 1.268548
Validation loss decreased (1.339883 --> 1.268548). Saving model ...
Epoch: 12     Training Loss: 2.590256      Validation Loss: 1.402471
Epoch: 13     Training Loss: 2.657677      Validation Loss: 1.310161
Epoch: 14     Training Loss: 2.602963      Validation Loss: 1.342396
Epoch: 15     Training Loss: 2.629120      Validation Loss: 1.221816
Validation loss decreased (1.268548 --> 1.221816). Saving model ...
Epoch: 16     Training Loss: 2.583704      Validation Loss: 1.268472
Epoch: 17     Training Loss: 2.531538      Validation Loss: 1.282012
Epoch: 18     Training Loss: 2.561271      Validation Loss: 1.274829
Epoch: 19     Training Loss: 2.605545      Validation Loss: 1.216648
Validation loss decreased (1.221816 --> 1.216648). Saving model ...
Epoch: 20     Training Loss: 2.532186      Validation Loss: 1.308733

```

## 2.0.5 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```

In [31]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    with torch.no_grad():
        for batch_idx, (data, target) in enumerate(loaders['test']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            # forward pass: compute predicted outputs by passing inputs to the model
            output = model(data)

```

```

    # calculate the loss
    loss = criterion(output, target)
    # update average test loss
    test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
    # convert output probabilities to predicted class
    pred = output.data.max(1, keepdim=True)[1]
    # compare predictions to true label
    correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
    total += data.size(0)

print('Test Loss: {:.6f}\n'.format(test_loss))

print('\nTest Accuracy: %2d%% (%2d/%2d)' % (100. * correct / total, correct, total))

```

In [32]: test(loaders\_transfer, model\_transfer, criterion\_transfer, use\_cuda)

Test Loss: 1.275388

Test Accuracy: 62% (522/836)

## 2.0.6 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```

In [35]: from PIL import Image
import matplotlib.pyplot as plt
%matplotlib inline

```

```

In [36]: ### TODO: Write a function that takes a path to an image as input
### and returns the dog breed that is predicted by the model.

```

```

# list of class names by index, i.e. a name can be accessed like class_names[0]
class_names = [item[4:].replace("_", " ") for item in train_data.classes]

def preprocess_img(img_path):
    img = Image.open(img_path)
    transform_pipeline = transforms.Compose([transforms.Resize((224,224)),
                                             transforms.ToTensor(),
                                             transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                                    std=[0.229, 0.224, 0.225])])

    img = transform_pipeline(img)
    img = img.unsqueeze(0)
    img = img.to('cuda')
    return img

```

```

hello, human!
0
200
400
600
800
1000
1200
1400
0 500 1000
You look like a ...
Chinese_shar-pei

```



Sample Human Output

```

def predict_breed_transfer(img_path):
    # load the image and return the predicted breed
    img = preprocess_img(img_path)
    model_transfer.eval()
    prediction = model_transfer(img)
    prediction = prediction.to('cpu')
    prediction = prediction.data.numpy().argmax()
    return class_names[prediction]

```

#### ## Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

### 2.0.7 (IMPLEMENTATION) Write your Algorithm

In [37]: *### TODO: Write your algorithm.*

*### Feel free to use as many code cells as needed.*

```

def show_image(img_path):
    img = Image.open(img_path)
    imgplot = plt.imshow(img)
    plt.show()

def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    if dog_detector(img_path):
        print("\nHello, dog! Your predicted breed is: " + predict_breed_transfer(img_path))

```

```

        print("\n")
        show_image(img_path)
    elif face_detector(img_path):
        print("\nHello, human! You look like a " + predict_breed_transfer(img_path))
        print("\n")
        show_image(img_path)
    else:
        print("\nyou are neither human nor dog!\n")
        show_image(img_path)

```

---

### ## Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

## 2.0.8 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

**Question 6:** Is the output better than you expected :) ? Or worse :( ? Provide at least three possible points of improvement for your algorithm.

**Answer:** My algorithm performs pretty well n the limited set of 6 images. It's able to correctly identify humans and also makes a correct prediction for the dog breed. However since the overall accuracy of the model is ~62% for the dog breed classifier, I can think of some possible improvements to increase accuracy of the model: 1. Try adding a couple of FC layers to the classifier section of the model. Might improve accuracy 2. Try using a deeper network architecture like ResNet instead of VGG16 2. Play around with different learning rate and optimization criteria like SGD, Adam, RMSprop etc to improve accuracy

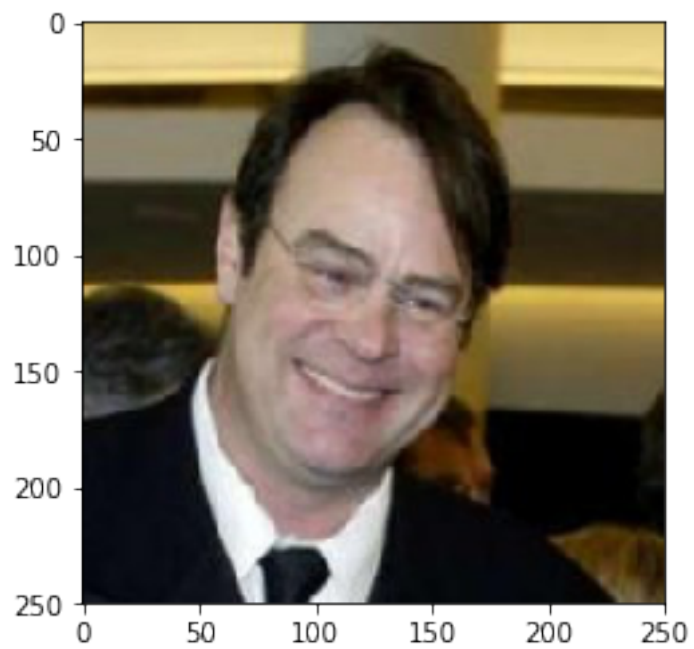
```

In [48]: ## TODO: Execute your algorithm from Step 6 on
        ## at least 6 images on your computer.
        ## Feel free to use as many code cells as needed.


        ## suggested code, below
        for file in np.hstack((human_files[:3], dog_files[:3])):
            run_app(file)

```

Hello, human! You look like a American staffordshire terrier




Hello, human! You look like a Greyhound







Hello, human! You look like a Cane corso




Hello, dog! Your predicted breed is: Bullmastiff



Hello, dog! Your predicted breed is: Bullmastiff



Hello, dog! Your predicted breed is: Bullmastiff



In [ ]: